

Artificial Neural Networks in Feature Extraction: A Review

Joaquim Leitão*
*jpleitao@dei.uc.pt

CISUC, Department of Informatics Engineering, Univesity of Coimbra, Portugal

Abstract—Lol, the abstract

Index Terms—Feature Extraction, Artificial Neural Networks.

1 INTRODUCTION

INTRO

Fazer intro curta

Dizer que foco na utilização de redes neuronais para efeitos de feature extraction (nao tanto de selection)

Colocar sigla para Feature Extraction - FE

4 pages

Do not forget to provide an overview of the paper!

2 BACKGROUND

The current section addresses the importance of Feature Extraction in Intelligent Systems. In the literature this concept is often presented in the sequel of the well-known curse of dimensionality phenomenon, as a collection of methods that allow a reduction in the problem's dimensionality.

2.1 Curse of Dimensionality

The curse of dimensionality refers to how certain learning algorithms may perform poorly when confronted with high-dimensional data. In many Pattern Recognition (PR) problems, low-dimension feature spaces may not allow for a good separation of the data, leading to low-accuracy classifications. By considering more features, a feature space where an hyperplane that perfectly separates the (training) data can be determined with a resulting increase in classification accuracy.

Nonetheless it has been demonstrated that, as dimensionality increases, the amount of training data needed to accurately generalise the produced classifier grows exponentially. In addition, a small *training samples-to-features* ratio may also degrade performance [1].

A practical implication of the curse of dimensionality is that a system designer should select a small number of features when confronted with a limited training dataset. A general accepted practice is to keep the number of training samples at least ten times higher than the dimensionality: $n/d > 10$ [1].

2.2 Dimensionality Reduction

In light of the discussion conducted in 2.1, a common interest in the vast majority of MPR applications is to keep the number of features as small as possible.

Feature Extraction and Feature Selection techniques are the most cited and recognised methods for achieving dimensionality reduction in PR problems. Feature Extraction algorithms propose a new set of features by performing transformations on a given feature set. On their turn, Feature Selection algorithms (ideally) select the best subset of features a provided feature set.

As pointed out by Jain *et al.* [1], feature extraction is typically applied before feature selection¹.

The main issues in dimensionality reduction, namely in feature extraction and selection, are related with the choice of criterion function and the appropriate dimension for the reduced feature space.

3 NEURAL NETWORKS IN FEATURE EXTRACTION

Artificial Neural Networks (ANNs) have been extensively applied in the field of machine learning, in both classification and regression tasks. Because of their self-learning and trained characteristics, ANNs have also been successfully applied in FE tasks.

Considering the simplest form of ANNs, feedforward networks, FE can be obtained in the output of each hidden layer: the application of the neurons' activation function to the layer's input produces an output that can be interpreted as a set of new features [1]. Depending on the neurons' activation function this new set of features can be a linear or nonlinear combinations of the features in the input feature set (usually is nonlinear).

In the remainder of the current section the most popular and referenced ANN architectures applied in FE tasks are intended to be covered. Such architectures include *Feedforward* (namely *Stacked Autoencoders*), *Convolutional*, *Recurrent*, *Radial Basis*, *Restricted Boltzmann Machines* and *Self-Organising Maps* [1]–[6].

FIXME: Rever a lista de arquiteturas para confirmar que não me esqueci de nenhuma!!

3.1 Feedforward

A feedforward ANN is one of the simplest examples of a neural network. The network is composed of several layers of neurons, where neurons in each layer are connected to neurons in the next layer: that is, connections between the units do not form a cycle. In the beginning of the current section feedforward networks were used as a motivation for the application of ANNs in FE tasks.

In its simplest forms, FE can be achieved through single or multi-layer feedforward ANNs by training the network to reconstruct the input signal in a conventional way, using a backpropagation algorithm (or other popular alternatives). This form of ANN training is considered unsupervised, as only unlabelled data is being used.

1. The idea is to start by defining a new feature set based on the original one, and then select the features that, hopefully, allow a complete separation of the training data.

In an opposite view, a multi-layer feedforward network may be trained to perform a classification task. As multiple layers are being considered, the output of the last hidden layer (which will provide the input to the output layer) can be considered as the extracted features from the provided input sample. Networks trained with this approach have been proven to suffer from well-known *vanishing gradient* problems [7].

In both scenarios the FE process is learned automatically in the early layers of the final network. In recent years, namely with advances in the field of deep learning, approaches of the first type tend to be more popular and widely used.

In this sense, training of feedforward networks for FE purposes usually comprises two steps: *Pre-training* and *Fine tuning*.

During pre-training a series of hidden layers are trained independently in an unsupervised way: Groups of 3-layer feed-forward nets (one for each hidden layer of the final network) are trained to reconstruct their respective inputs at the output layer. The idea is to make each network learn an encoder and a decoder². In this process, the first network is trained with the input data; the second network is trained with the output of the first network's encoder, and so on. This process is illustrated in figure 1.

When pre-training is complete, the encoders are stacked together. Due to this last step, ANNs developed in this way are often referred to as *Stacked autoencoders*.




Figure 1. Stacking autoencoders. Taken from [8].

After the described pre-training, a fine tuning step is commonly carried out by considering an additional layer (usually with a softmax activation function) and training the network with backpropagation (or other algorithm), to lightly adjust the weights of the pre-trained layers.

At this moment it is important to address one point in the discussion: based on what has been presented, autoencoders may simply learn the identity function for both the encoder and the decoder, which is not desired. To prevent this, certain activation functions can be forced, producing *regularized autoencoders* such as sparse autoencoders [9], denoising autoencoders [10] and contractive autoencoders [11].

Feedforward networks, namely *Stacked autoencoders*, can be applied for Feature Extraction purposes in a series of applications, including image processing and classification and speech recognition [9], [12], just to name a few.

3.2 Convolutional Neural Networks

Convolutional Neural Networks (CNN) are ANNs designed to process data that comes in the form of multiple arrays, such as an image of a sequence of 1D signals. CNNs are based on four key ideas that take advantage of the properties of natural signals: local connections, shared weights, pooling and multi-layer usage [13]. Similarly to the feedforward case, CNNs

can be used to automatically extract features and perform classification tasks.

The architecture of a CNN comprises several stages: In the first few two types of layers are alternated: *convolutional* and *pooling* - In the literature it is common to refer to a pair of convolutional and pooling as one hidden layer of the CNN, or a *slice*. Following these sequences of alternating convolutional and pooling layers, a fully connected layer is considered in order to produce the network's output. Figure 2 presents a conceptual example of a CNN.




Figure 2. Conceptual example of a Convolutional Neural Network. Taken from [14].

In short, units in convolutional layers are organised in feature maps, where for each map a subset of units from one hidden layer serves as input to the next hidden layer. Like in other ANN architectures, each input sample has an associated weight. The same weight vector must be applied to the neurons in a single depth slice.

The logic behind the convolutional layer is to perform, for each map in a given hidden layer, a convolution of the neurons' weights with the input volume. The result of this local weighted sum is then passed through a non-linear activation function, such as a *ReLU*³.

Because of the mentioned weight sharing principle, it is common to refer to the sets of weights as a *filter* or *kernel* that is applied to the input. For this reason, the convolutional layer is often described as consisting in the application of several filters to the input (resulting from the different feature maps considered). It can also be said that the role of this layer is to detect local conjunctions of features [13].

In a different view, units in the pooling layers enforce a form of non-linear down-sampling, by seeking to merge semantically similar features into one [13]. Several non-linear functions have been proposed to this purpose, a typical pooling function is the *max pooling*, which computes the maximum of a local set of units in one (or few) feature map.

From a feature extraction and dimensionality reduction point of view, the described sequence of convolutional and pooling layers is responsible for reducing the dimensionality of the provided input set. In some scenarios a normalisation layer may be considered; however, as pointed out in [15], this practice has fallen out of favor because its contribution has been shown to be minimal.

In the vast majority of the scenarios, CNNs are applied in classification problems, and for that reason additional fully-connected layers are considered where an high-level reasoning is performed. Regarding the network's training, the application of a backpropagation gradient through a CNN is similar to regular deep networks, allowing all the weights in all the layers and filters to be trained.

Convolutional Neural Networks have been extensively applied for Feature Extraction purposes, with the most sounding

2. The encoder is learned from the input to the hidden layer, while the decoder is learned from the hidden to the output layer.

3. Rectified Linear Unit.

applications in fields such as image [16], [17], video [18] and speech processing and recognition [19], [20], natural language processing [21] and analysis of other sequences of 1D signals. CNNs were also applied to develop an artificial intelligence for the board game Go [22], [23]

3.3 Recurrent Neural Networks

Recurrent Neural Networks (RNN) are ANNs designed to process sequential information. Unlike in Feedforward networks, where each layer can only be connected to the next one, forming a directed acyclic graph, connections in a RNN can form directed cycles. That is, in a RNN one neuron layer can be connected to a previous layer, or itself. Figure 3 presents a RNN with one hidden layer.




Figure 3. Illustration of a Recurrent Neural Network. Taken from [24].

Two distinct types of recurrence can be considered in RNN: If the outputs of a given layer are connected to the inputs of the same layer, then *internal recurrence* is being considered; however, if the output layer is connected to the input layer, *external recurrence* is being considered. Current RNN architectures tend to consider *internal recurrence* rather than *external recurrence*.

RNNs can be seen as structures with *memory*, allowing them to capture information about what previous computations. This is the main reason why RNNs are so popular and effective when dealing with sequential data, such as time series. In theory RNNs can make use of information in arbitrarily long sequences, but in practice they are limited to looking back only a few steps.

Similarly to Feedforward networks, several activation functions can be considered in the neurons of a RNN, of which sigmoid and softmax activation functions are examples.

Studying the literature in the topic of RNNs several architectures for RNNs can be found. the most popular architectures proposed in the literature include *Elman* [25], [26] and *Jordan* [27] networks. Because of space restrictions, a detailed description of the different structures will not be presented. The reader is directed to the cited works for a more detailed coverage.

With respect to the training of this networks, *Backpropagation Through Time* (BPTT) is among the most referenced and known algorithms applied to RNN. BPTT is defined as an adaptation of the popular backpropagation⁴ training method applied in feed-forward neural networks. This method proposes an “*unfolding*” of the recurrences in the network, as presented in figure 4.

Applying this technique, a feedforward neural network is obtained, and the backpropagation algorithm can be applied. During the training process the BPTT algorithm progressively unfolds the network, leading to a continuous increase in the size of the network. In theory, BPTT does not define a limit for the unfolding, although approaches such as *Truncated BPTT* have been proposed to achieve this by defining a fixed number of layers in the network.

Other alternatives to these methods can also be identified in the literature, of which *Kalman Filters* [28], *Real-Time Recurrent Learning* (RTRL) [29] and *Long Short-Term Memory* (LSTM) [30].

4. A detailed explanation on how the backpropagation method works can be found in [7].




Figure 4. Illustration of a Recurrent Neural Network. Taken from [24].

The main feature of the RTRL algorithm has to do with its mathematical formulation: the derivatives are explicitly computed, and therefore, no unfolding is performed. By its turn, LSTM avoids the well-known vanishing gradient problems and introduces a distinct architecture, where units are implemented in blocks containing gates used to control the flow of information into or out of their memory. Again, due to size restrictions on the current paper, the reader is deferred to the cited literature for a more in-depth coverage of these techniques.

Recurrent Neural Networks have been extensively applied for Feature Extraction purposes, with the most sounding applications in fields such as image classification [31]; time series prediction, namely in EEG signal prediction [32], speech recognition [33], among others. The work of Lee and Tashev [33], focusing on speech recognition, applies RNNs in automatic high-level feature representation, which can be considered as a FE task. Indeed, a transformation of raw input data to a representation that can be effectively exploited in machine learning tasks is carried out by a RNN in this work.

3.4 Radial Basis Functions

3.5 Restricted Boltzmann Machines

3.6 Self-Organising Maps

In Self-Organising Maps (SOM), another type of ANN used for NLFE [1], [34], a low-dimension discretized representation of the input space is produced. Such representation is called a map.

A SOM consists of neurons with associated weight vectors (with the same dimension as the input). The neurons are positioned in the map space, usually in a two-dimensional regular spacing in either a hexagonal or rectangular grid (instead of being organised in layers, like in feedforward ANNs).

During the training of a SOM, inputs are presented to the network and, for each, the weight vector closest to the input vector is identified. Then, the weights of all the neurons in the neighbourhood of the winner neuron (that is, the one whose weight vector was selected) are updated, so that they move towards the input vector. In this way, weight vectors of neighbouring neurons in the grid are likely to represent inputs which are close in the original feature space - a “topology-preserving” map is, thus, formed [35], [36]. An example of an application where SOM was used for FE can be found in [16], although many more are available in the literature.

4 CONCLUSION

Lol, conclusion

REFERENCES

- [1] A. K. Jain, R. P. W. Duin, and J. Mao, “Statistical pattern recognition: A review,” *IEEE Transactions on pattern analysis and machine intelligence*, vol. 22, no. 1, pp. 4–37, 2000.

- [2] J. Masci, U. Meier, D. Cireşan, and J. Schmidhuber, "Stacked convolutional auto-encoders for hierarchical feature extraction," in *International Conference on Artificial Neural Networks*, pp. 52–59, Springer, 2011.
- [3] G. Kvascev, M. G. Kvascev, and Z. Djurovic, "Radial basis function network based feature extraction for improvement the procedure of sourcing neolithic ceramics," in *Neural Network Applications in Electrical Engineering (NEUREL), 2012 11th Symposium on*, pp. 95–100, IEEE, 2012.
- [4] O. Fabius and J. R. van Amersfoort, "Variational recurrent auto-encoders," *arXiv preprint arXiv:1412.6581*, 2014.
- [5] K. Cho, B. Van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, "Learning phrase representations using rnn encoder-decoder for statistical machine translation," *arXiv preprint arXiv:1406.1078*, 2014.
- [6] E. Marchi, F. Vesperini, S. Squartini, and B. Schuller, "Learning phrase representations using rnn encoder-decoder for statistical machine translation," *Computational Intelligence and Neuroscience*, vol. 2017, 2017.
- [7] M. A. Nielsen, *Neural Networks and Deep Learning*. Determination Press, 2015.
- [8] R. Barata, "Stacked autoencoders." Lecture. Advanced Topics in Cognitive Models. University of Coimbra, 2017.
- [9] P. Vincent, H. Larochelle, I. Lajoie, Y. Bengio, and P.-A. Manzagol, "Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion," *Journal of Machine Learning Research*, vol. 11, no. Dec, pp. 3371–3408, 2010.
- [10] J. Ngiam, A. Coates, A. Lahiri, B. Prochnow, Q. V. Le, and A. Y. Ng, "On optimization methods for deep learning," in *Proceedings of the 28th International Conference on Machine Learning (ICML-11)*, pp. 265–272, 2011.
- [11] S. Rifai, P. Vincent, X. Muller, X. Glorot, and Y. Bengio, "Contractive auto-encoders: Explicit invariance during feature extraction," in *Proceedings of the 28th international conference on machine learning (ICML-11)*, pp. 833–840, 2011.
- [12] X. Lu, Y. Tsao, S. Matsuda, and C. Hori, "Speech enhancement based on deep denoising autoencoder," in *Interspeech*, pp. 436–440, 2013.
- [13] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, no. 7553, pp. 436–444, 2015.
- [14] I. Arel, D. C. Rose, and T. P. Karnowski, "Deep machine learning-a new frontier in artificial intelligence research [research frontier]," *IEEE Computational Intelligence Magazine*, vol. 5, no. 4, pp. 13–18, 2010.
- [15] "Cs231n convolutional neural networks for visual recognition." <http://cs231n.github.io/convolutional-networks/>. Accessed: 2017-04-10.
- [16] S. Lawrence, C. L. Giles, A. C. Tsoi, and A. D. Back, "Face recognition: A convolutional neural-network approach," *IEEE transactions on neural networks*, vol. 8, no. 1, pp. 98–113, 1997.
- [17] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "Imagenet classification with deep convolutional neural networks," in *Advances in neural information processing systems*, pp. 1097–1105, 2012.
- [18] A. Karpathy, G. Toderici, S. Shetty, T. Leung, R. Sukthankar, and L. Fei-Fei, "Large-scale video classification with convolutional neural networks," in *Proceedings of the IEEE conference on Computer Vision and Pattern Recognition*, pp. 1725–1732, 2014.
- [19] Y. LeCun, Y. Bengio, et al., "Convolutional networks for images, speech, and time series," *The handbook of brain theory and neural networks*, vol. 3361, no. 10, p. 1995, 1995.
- [20] O. Abdel-Hamid, A.-r. Mohamed, H. Jiang, and G. Penn, "Applying convolutional neural networks concepts to hybrid nn-hmm model for speech recognition," in *Acoustics, Speech and Signal Processing (ICASSP), 2012 IEEE International Conference on*, pp. 4277–4280, IEEE, 2012.
- [21] B. Hu, Z. Lu, H. Li, and Q. Chen, "Convolutional neural network architectures for matching natural language sentences," in *Advances in neural information processing systems*, pp. 2042–2050, 2014.
- [22] C. Clark and A. Storkey, "Teaching deep convolutional neural networks to play go," *arXiv preprint arXiv:1412.3409*, 2014.
- [23] C. J. Maddison, A. Huang, I. Sutskever, and D. Silver, "Move evaluation in go using deep convolutional neural networks," *arXiv preprint arXiv:1412.6564*, 2014.
- [24] "Recurrent neural networks tutorial, part 1 – introduction to rnns." <http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-1-introduction-to-rnns/>. Accessed: 2017-04-11.
- [25] J. L. Elman, "Finding structure in time," *Cognitive science*, vol. 14, no. 2, pp. 179–211, 1990.
- [26] J. L. Elman, "Distributed representations, simple recurrent networks, and grammatical structure," *Machine learning*, vol. 7, no. 2–3, pp. 195–225, 1991.
- [27] M. I. Jordan, "Serial order: A parallel distributed processing approach," *Advances in psychology*, vol. 121, pp. 471–495, 1997.
- [28] R. E. Kalman et al., "A new approach to linear filtering and prediction problems," *Journal of basic Engineering*, vol. 82, no. 1, pp. 35–45, 1960.
- [29] R. J. Williams and D. Zipser, "A learning algorithm for continually running fully recurrent neural networks," *Neural computation*, vol. 1, no. 2, pp. 270–280, 1989.
- [30] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [31] A. Graves and J. Schmidhuber, "Offline handwriting recognition with multidimensional recurrent neural networks," in *Advances in neural information processing systems*, pp. 545–552, 2009.
- [32] S. C. Prasad and P. Prasad, "Deep recurrent neural networks for time series prediction," *arXiv preprint arXiv:1407.5949*, 2014.
- [33] J. Lee and I. Tashev, "High-level feature representation using recurrent neural network for speech emotion recognition," in *Interspeech 2015, ISCA - International Speech Communication Association*, September 2015.
- [34] Z. M. Hira and D. F. Gillies, "A review of feature selection and feature extraction methods applied on microarray data," *Advances in bioinformatics*, vol. 2015, 2015.
- [35] T. Kohonen, "Self-organising maps," *Springer Series in Information Science*, vol. 30, Berlin, 1995.
- [36] T. Villmann, R. Der, M. Herrmann, and T. M. Martinetz, "Topology preservation in self-organizing feature maps: exact definition and measurement," *IEEE Transactions on Neural Networks*, vol. 8, no. 2, pp. 256–266, 1997.